

Assignment 2: Gradient Descent for Linear and Logistic Regression

Loyola Marymount University

Professor Alex Wong

CSMI-533: Data Science and Machine Learning

Implement logistic regression and linear regression using gradient descent

There are many ways to train logistic and linear regression models. For this assignment we will be using gradient descent

We will implement the `GradientDescentOptimizer` class using the Gradient Descent Algorithm. We will also implement the `LinearRegressionGradientDescent` class and the `LogisticRegressionGradientDescent` class -- both of which will use the `GradientDescentOptimizer` class to update their respective weights.

We will apply the logistic regression models to the Wisconsin breast cancer dataset and the digits dataset (for classifying 7s and 9s).

We will apply the linear regression models to the Boston housing price dataset and the diabetes dataset.

A skeleton of the classes is provided in `assignment2.py`

You will complete the following functions for the assignment:

In the `GradientDescentOptimizer` class:

- 1) `__compute_gradients` : Returns the gradient of the logistic, mean squared or half mean squared loss
- 2) `update` : Updates the weight vector based on logistic, mean squared or half mean squared loss

In the `LogisticRegressionGradientDescent` class:

- 1) `fit` : Fits the model to x and y by updating the weight vector using gradient descent
- 2) `predict` : Predicts a label for each feature vector x
- 3) `score` : Predicts labels based on feature vector x and computes the mean accuracy of the predictions

In the `LinearRegressionGradientDescent` class:

- 4) `fit` : Fits the model to x and y by updating the weight vector using gradient descent
- 5) `predict` : Predicts a real value for each feature vector x
- 6) `score` : Predicts labels based on feature vector x and computes the mean squared error

Here is the sample output:

Results on Wisconsin breast cancer dataset using scikit-learn Logistic Regression model

Training set mean accuracy: 0.9590

Testing set mean accuracy: 0.9649

Results on digits 7 and 9 dataset using scikit-learn Logistic Regression model

Training set mean accuracy: 1.0000

Testing set mean accuracy: 0.9722

Results on Boston housing price dataset using scikit-learn Linear Regression model

Training set mean accuracy: 23.2335

Testing set mean accuracy: 10.8062

Results on diabetes dataset using scikit-learn Linear Regression model

Training set mean accuracy: 2991.9850

Testing set mean accuracy: 1735.9381

Results on Wisconsin breast cancer dataset using our Logistic Regression model

Training set mean accuracy: 0.0000

Testing set mean accuracy: 0.0000

Results on digits 7 and 9 dataset using our Logistic Regression model

Training set mean accuracy: 0.0000

Testing set mean accuracy: 0.0000

Results on Boston housing price dataset using our Linear Regression model

Training set mean accuracy: 617.1287

Testing set mean accuracy: 369.2698

Results on diabetes dataset using our Linear Regression model

Training set mean accuracy: 29088.9673

Testing set mean accuracy: 28946.6889

Submission:

You will submit the following to Bright Space

- 1) <last_name>_<first_name>_assignment2.py

Grading:

I will be executing the assignment using the following command:

```
python <last_name>_<first_name>_assignment2.py
```

Your code must run for me to assign points!

Your assignment will be graded on:

- 1) the correctness of your implementation of the GradientDescentOptimizer class
- 2) the correctness of your implementation of the LogisticRegressionGradientDescent class
- 3) the correctness of your implementation of the LinearRegressionGradientDescent class
- 4) results of your LogisticRegressionGradientDescent class on the Wisconsin breast cancer dataset and the digits dataset
- 5) results of your LinearRegressionGradientDescent class on the Boston housing price dataset and the diabetes dataset

Note: your results should deviate no more than 5%-7% of the performance given by scikit-learn

Late Policy:

For each day the assignment is late, 50% of its worth will be deducted, e.g. 100% on time, 50% 1 day late, 25% 2 days late, etc.